IT IS SLOW

Performance analysis and optimisation on the JVM



Materials

Code

git clone https://github.com/srvaroa/jbcnconf

check README for build/exec instructions

Apps

JDK (Hotspot > OpenJDK, but both fine)

Vagrant + VirtualBox (optional)

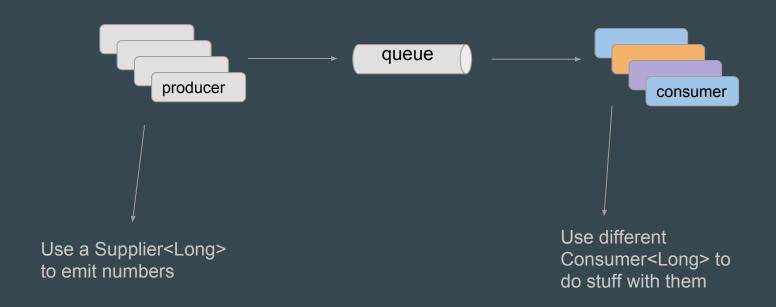
Slides from BcnJUG meetup

http://bit.ly/1UJSeTW

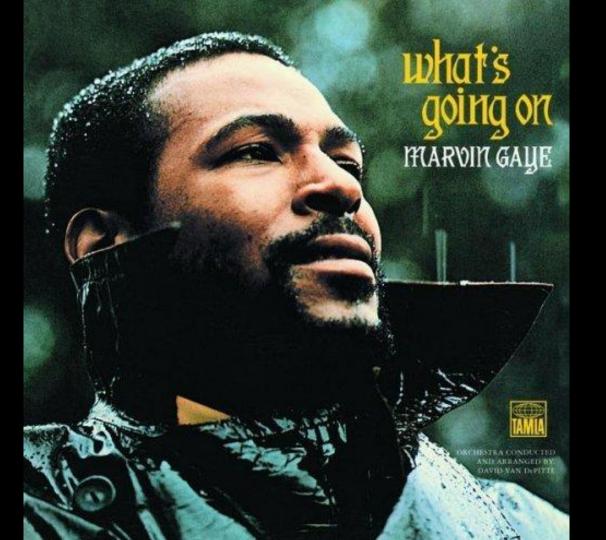
Agenda

- Learn to analyze aspects of an app's performance
- Tools
- Some common pathologies
- Some JDK tools and patterns to address them

SimpleApp



Concurrency



Why?

Usually, only available tool (e.g.: remote / production / customer deployments)

CPU

```
jstack -l $PID 1000
```

Memory

```
jstat -gc $PID $INTERVAL

jmap -heap $PID

jmap -histo $PID

jmap -histo:live $PID
```

Performance counters

jcmd \$PID PerfCounter.print

Flight Recorder + Mission Control

OSX: /Library/Java/Home/bin/jmc

Linux: \$JAVA HOME/bin/jmc

Run App

Attach MC

Flight Recorder + Mission Control

```
-XX:+UnlockCommercialFeatures
-XX:+FlightRecorder
-XX:StartFlightRecording=duration=60s, filename=rec.jfr, settings=profile

Profiles (editable) $JAVA_HOME/jre/lib/jfr/
-XX:FlightRecorderOptions=loglevel=trace
```

Experiments

Change queue sizes

Change producer + consumer threads (symmetric, asymmetric)

Change supplier implementation

Comment out one of the Consumer implementations

Use different logging mechanisms

Benchmarking

JMH

"JMH is a Java harness for building, running, and analysing nano/micro/milli/macro benchmarks written in Java and other languages targeting the JVM."

http://openjdk.java.net/projects/code-tools/jmh/

Hello World:

com.github.srvaroa.jmh.IterableBenchmark

Documentation:

https://github.com/midonet/midonet/blob/master/docs/micro-benchmarks.md

Observations

Observations

Contention is expensive

Easy synchronization tools are too inefficient

• synchronized is borderline harmful

Better sync tools requires complex code

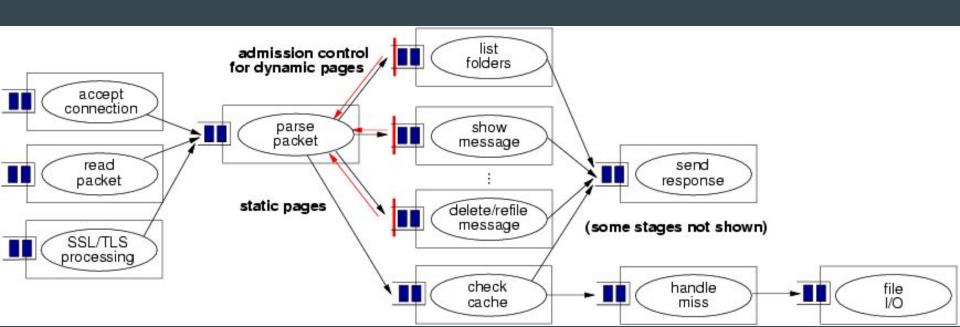
Optimisation requires hardware-specific knowledge

Lots of black boxes underneath (JDK, JVM, OS, CPUs, Network cards...)

Queues are troublesome at scale

Queues are key in SEDA architecture, but a major hotspot

https://issues.apache.org/jira/browse/CASSANDRA-10989: Move from SEDA to TPC (thread per core)



OneToOneApp

Paired producer - consumer

Use JcToolsSpscBackedQueue and check code inside

Frameworks

ReactiveX

https://github.com/ReactiveX/RxJava/ - in provided VM

Rx is push based. Paradigm change: work enters a pipeline, functions in a single thread until told otherwise.

Does this make sense in light of what we saw above?

Mental model change: from stages communicated with queues to flows of data

Caveat: not all Rx operators are optimal.

Akka (Actors + Streams)

https://github.com/akka/akka - in provided VM

Good points

Again, avoid contention in processing units (Actors)

Bad points

Queues, CPU hopping

Small units of work have large contention on Dispatcher queues

Notes on RX & Actors

They are both a means to model your system, let you define separately:

Stages, units of work (Akka actors)

Flows of computation (Rx, Akka streams)

How they are scheduled (Akka: put a group of actors on a particular dispatcher)

Great in many applications, but some pitfalls, tradeoffs:

Example: sending work to another dispatcher can be expensive (wake up thread)

Always tune / profile

Perf & Flame Graphs

Perf (GNU/Linux)

Setup VM

```
$ cd vm
$ vagrant up
```

By hand: check end of vm/VagrantFile for packages

Build JAR locally, then copy into VM:

```
$ cp ./build/libs/jbcnconf-0.1-SNAPSHOT-all.jar ../vm/data # from code/
$ cd ../vm && vagrant ssh
```

File will be at /vagrant_data/jbcnconf-0.1-SNAPSHOT-all.jar

Perf stat (not available in VM)

```
perf stat -d -p $PID # also: cat /proc/$PID/status
```

```
Performance counter stats for process id '46185':
```

```
46802,773132 task-clock
                                           0,367 CPUs utilized
       24521 context-switches
                                           0.001 M/sec
       1056 CPU-migrations
                                           0,000 M/sec
        450 page-faults
                                           0,000 M/sec
129724461306 cycles
                                           2,772 GHz
                                                                         [40,03%]
120451823606 stalled-cycles-frontend
                                          92,85% frontend cycles idle
                                                                         [40,40%]
101327998695 stalled-cycles-backend
                                         78,11% backend cycles idle
                                                                          [40.56%]
  5710255927 instructions
                                           0.04 insns per cycle
                                          21,09 stalled cycles per insn [50,69%]
  1334448420 branches
                                          28,512 M/sec
                                                                         [50,80%]
                                         2,61% of all branches
   34780770 branch-misses
                                                                          [50,69%]
  1540840733 L1-dcache-loads
                                          32,922 M/sec
                                                                         [50,22%]
  4551850342 L1-dcache-load-misses
                                         295,41% of all L1-dcache hits
                                                                          [50,02%]
    71937736 LLC-loads
                                           1.537 M/sec
                                                                          [39,83%]
    42543914 LLC-load-misses
                                           59,14% of all LL-cache hits
                                                                         [39,78%]
```

127,428135422 seconds time elapsed

Flame Graphs (available in VM)

Visualization for sampled stack frames. Great to see where your CPU is spending time.

http://www.brendangregg.com/FlameGraphs/cpuflamegraphs.html

Java: stacks showed only application & JDK code paths, not OS.

Since 1.8 u60 b19: -XX: +PreserveFramePointer enables end-to-end stacks

http://techblog.netflix.com/2015/07/java-in-flames.html

http://www.brendangregg.com/FlameGraphs/cpu-mixedmode-flamegraph-java.svg

Flame Graphs

On Java, config instructions (already in VM):

https://github.com/brendangregg/FlameGraph http://isuru-perera.blogspot.com.es/2015/07/java-cpu-flame-graphs.html

Flame Graphs

Preconfigured in provided VM, easy to generate:

```
$ java -cp $JAR -XX:+PreserveFramePointer com.github.srvaroa.queue.SimpleApp
$ jps -m
11633 Jps -m
11606 SimpleApp
$ cd perf-map-agent $$ bin/perf-java-flames 11606
Recording events for 15 seconds (adapt by setting PERF RECORD SECONDS)
[ perf record: Woken up 1 times to write data ]
[ perf record: Captured and wrote 0.214 MB /tmp/perf-11606.data (~9357 samples)
Flame graph SVG written to PERF FLAME OUTPUT='/home/vagrant/perf-map-
agent/flamegraph-11606.svg'.
$ cp /vagrant/perf-map-agent/flamegraph-11606.svg /vagrant data
```

Use browser to open in host at ibanconf/ym/data/flamegraph-11606 syg



Q & A

Thanks!

Galo Navarro @srvaroa

Guillermo Ontañón *@gontanon*